

Solutions for Homework Assignment #5

Answer to Question 1. CLAIM: $\frac{5}{2}$ is the best (i.e., smallest) upper bound for $T(n)/n$ in the list L .

- **First show for all n :** $T(n)/n \leq \frac{5}{2}$.

To show this, we use the accounting method, and charge operations as follows:

(a) $\frac{5}{2}$ for each `INSERT(x)` operation, and (b) 0 for each `OUTPUTANDREDUCE()`.

We claim that the following *credit invariant* holds:

The total credit of any array A of size n is at least $n \times \frac{3}{2}$ (*)

Suppose (*) holds just before an operation o on an array A of size n , we need to show that (*) also holds immediately after operation o . There are two cases:

(a) o is an insert operation. Since (*) holds, before o is executed the array A has at least $n \times \frac{3}{2}$ credits. After executing o , the array's amount of credits is at least $n \times \frac{3}{2} + (\frac{5}{2} - 1) = (n+1) \times \frac{3}{2}$, and the array's size is $n+1$. So (*) still holds.

(b) o is a `OUTPUTANDREDUCE()` operation. Before o , the array A has at least $n \times \frac{3}{2}$ credits. Executing o costs exactly n credits (because o prints the n elements of A), so array A has at least $n \times \frac{3}{2} - n = n \times \frac{1}{2} = \frac{n}{3} \times \frac{3}{2}$ credits left.

Since the size of A is now at most $n' = n/3$, A has at least $n' \times \frac{3}{2}$ credits, and therefore the invariant (*) still hold.

So (*) always holds. From (*), the amount of accumulated credits is always positive. So the total amount charged to any sequence of $n \geq 1$ operation (more than) covers the total cost of executing these operations. Since we charge at most $\frac{5}{2}$ per operation, we have $T(n)/n \leq \frac{5}{2}$.

ALTERNATIVE PROOF:

To show $T(n)/n \leq \frac{5}{2}$, we prove that if we charge:

(a) $\frac{5}{2}$ for each `INSERT(x)` operation, and (b) 0 for each `OUTPUTANDREDUCE()`

then we always cover the total cost of any sequence of operations.

Since the actual cost of doing an insert is exactly 1, after an element is inserted, it has a remaining credit of $\frac{5}{2} - 1 = \frac{3}{2}$. In fact, we will maintain the invariant that every element in the array always has $\frac{3}{2}$ credit on it.

To pay for the cost of executing a `OUTPUTANDREDUCE()` operation, we first take one credit from every element that is printed (so if this prints k elements, this takes k credits out of the array).

Note that after doing so, the remaining credit attached to each element is now $\frac{3}{2} - 1 = \frac{1}{2}$.

Since the `OUTPUTANDREDUCE()` operation effectively removes the last 2/3rd of the array (only the first third of the array remains), we now give to *each* element in the first third of the array the $2 \times \frac{1}{2}$ credits of *two* elements located in the removed part of the array.

After doing this, we are back to a state where each element in the array has $\frac{1}{2} + 1 = \frac{3}{2}$ credit attached to it.

- **Then show that for some n :** $T(n)/n > 2$.

Consider the following sequence of operations: first insert 27 elements, then print 27, print 9, print 3, and finally print 1.

This is a sequence of $n = 27 + 4 = 31$ operations whose total cost $T(n) = 27 + 27 + 9 + 3 + 1 = 67$.

So $T(n)/n = 67/31 > 2$.

Answer to Question 2.

a. Assume that the directed graph G is given by its adjacency matrix A , so $A[i, j] = 1$ if and only if G has an edge from node i to node j . To find out if the graph G has a supersource, the basic idea is to use each comparison “ $A[i, j] = 0$?”, which can be done in $O(1)$ -time using matrix A , to *eliminate one of node i or j* as a possible candidate for being a supersource: if $A[i, j] = 0$ (for $i \neq j$), then node i cannot be supersource; and if $A[i, j] \neq 0$, then node j cannot be a supersource. So with $n - 1$ questions, i.e., in $O(n)$ -time, we can eliminate $n - 1$ nodes, and we remain with only one possible supersource candidate, say node s . It is now easy to see whether s is indeed a supersource by checking whether the following two conditions hold: (a) $A[s, j] = 1$ for every $j \neq i$ (there is an edge from s to every other node), and (b) $A[j, s] = 0$ for every j (there are no edges into s). This check, which involves scanning a row and a column of A , takes $O(n)$ -time. More precisely, the pseudo-code of the algorithm is:

```

s ← 1           {s is the current supersource candidate}
j ← 2           {j is another candidate }
while j ≤ n do
    if A[s, j] = 0 then s ← j    { if there is no edge from s to j, then j becomes the new supersource candidate}
    j ← j + 1
{ At the exit of the above loop, if G has a supersource then it must be s }
{ We now check whether s is indeed a supersource.}
for j ← 1 to n
    if A[s, j] = 0 and j ≠ s then print “G does not have a supersource”; stop.
    if A[j, s] = 1 then print “G does not have a supersource”; stop.
print “G has a supersource and it is” s ; stop

```

PROOF OF THE ALGORITHM: It is not difficult to prove that at the end of each iteration of the **while** loop the following invariant holds:¹

INVARIANT: $[1 \leq s < j \leq n + 1]$ **and** [if G has a supersource **and** it is among $\{1, 2, \dots, j - 1\}$ **then** it must be s].

At the exit of the **while** loop $j = n + 1$. So, by the above loop invariant, if there is a supersource among $\{1, 2, \dots, n\}$ then it must be s . The code that follows the **while** loop simply checks whether s is indeed a supersource (i.e., whether there is an edge from s to every other node but there is no edge into s).

b. The **while** loop above executes $n - 1$ times, so it accesses the matrix A exactly $n - 1$ times. The **for** loop executes n times, and each iteration accesses the matrix A twice, so it accesses the matrix $2n$ times. So the total number of accesses to A of this algorithm is $3n - 1$.

c. Now assume that the graph G is given by its adjacency lists L . To check whether there is an edge from a node i to a node j now requires traversing the list $L[i]$ to see whether j is in this list. In the worst case this takes $\Theta(n)$ time. So the above algorithm now takes $\Theta(n^2)$ time in the worst-case.

A more efficient algorithm (for the case where G is given by its adjacency lists) is sketched below:

1. For each $i, 1 \leq i \leq n$, check whether the list $L[i]$ contains exactly $n - 1$ nodes. If there is *exactly one* node s such that $L[s]$ contains exactly $n - 1$ nodes, then s is a supersource candidate; otherwise print “ G does not have a supersource” and stop.

This step takes $\Theta(n + m)$ time in the worst-case.

2. If Part 1 above yielded a supersource candidate s , check that for all $i, 1 \leq i \leq n$, node s is *not* on list $L[i]$; if so print “ G has a supersource and it is” s , else print “ G does not have a supersource”, and stop.

This step also takes $\Theta(n + m)$ time in the worst-case.

Answer to Question 3.

a. We run a BFS starting from any node of G to create a BFS tree. Then we direct every tree edge from the parent to the child node. The non-tree edges are oriented arbitrarily. Because the running time of this algorithm

¹By convention, the end of the 0-th iteration of the loop is the beginning of the first iteration of the loop.

is dominated by the running time of BFS, the algorithm runs in time $O(m + n)$. We claim that, after the edges are oriented, every node of G , except the root of the BFS tree has at least one edge going into it. Indeed, every node in the BFS tree except the root has a parent node, and the edge connecting them is oriented from the parent to the child node.

b. Since each $v \in V$ has at least one edge (u, v) going into it (for some $u \in V$) in the oriented graph G , and every directed edge goes into exactly one vertex, the oriented graph G must have at least n directed edges. Since each undirected edge of the original graph yields exactly one directed edge in the oriented graph, the undirected graph must have at least n undirected edges as well.

c. We run BFS once, from an arbitrary vertex. Because the graph is connected and has at least n edges, while the BFS tree has $n - 1$ edges, there must be at least one non-tree edge; let us call this edge $e = (u, v)$. We direct this edge from v to u , and remove it from the graph to get a new graph $G' = (V, E - \{e\})$. Because G' still has all edges that were in the BFS tree, it is still connected. We run BFS starting from u , and orient all edges in the new BFS tree from parent to child node; we orient all other edges arbitrarily. By the analysis in the first subquestion, this leaves only vertex u possibly without any edge going into it. However, we oriented edge e so that it goes into u , and, thus, we have at least one edge going into every vertex. This proves the correctness of the algorithm. The running time is $O(m + n)$ because it is dominated by the two runs of BFS, each of them taking time $O(m + n)$.